



Vector Search

Evolution of Searching

- Exact matching
"Car" == "Car"
- Partial/Wildcard matching
"Car*" == "Carpet"
- Full-text matching
"My car is red" => ["my", "car", "red"] == "car"
- Semantic matching
"Signature red car" => "Ferrari"



Semantic similarity- how?

- Take Phrase1
- Convert it to NumericalRepresentation1
- Take Phrase2
- Convert it to NumericalRepresentation2
- Compute distance between NR1 and NR2
- Distance == Similarity



How to represent words numerically?

- “Beach Towel”
- beach 0.91
- sun 0.87
- water 0.78
- sand 0.84
- swimwear 0.72
- vacation 0.76
- sunscreen 0.70
- summer 0.86
- waves 0.75
- sunbathing 0.82
- surf 0.68
- pool 0.73
- heat 0.65
- resort 0.77
- winter 0.24
- snow 0.19
- wind 0.30
- skiing 0.16
- jacket 0.22



How to represent words numerically?

- “Beach Towel”

• beach	0.91
• sun	0.87
• water	0.78
• sand	0.84
• swimwear	0.72
• vacation	0.76
• sunscreen	0.70
• summer	0.86
• waves	0.75
• sunbathing	0.82
• surf	0.68
• pool	0.73
• heat	0.65
• resort	0.77
• winter	0.24
• snow	0.19
• wind	0.30
• skiing	0.16
• jacket	0.22

- “Sunglasses”

• beach	0.84
• sun	0.93
• water	0.72
• sand	0.75
• swimwear	0.70
• vacation	0.80
• sunscreen	0.78
• summer	0.89
• waves	0.70
• sunbathing	0.82
• surf	0.66
• pool	0.73
• heat	0.76
• resort	0.78
• winter	0.28
• snow	0.34
• wind	0.42
• skiing	0.38
• jacket	0.25



How to compute similarity?

1. Represent numbers as vectors - “vector embeddings”

sunglasses $\rightarrow \mathbf{v}_1 = [0.84, 0.93, \dots, 0.25]$


beach $\rightarrow \mathbf{v}_2 = [0.91, 0.87, \dots, 0.22]$

2. Compute “Cosine Similarity”

1.0 \rightarrow perfectly similar (same direction)

0.0 \rightarrow unrelated (orthogonal)

0.81 = 81%



How to automate this?

- AI, of course!
- “Text Embedding Models” with various “vector dimensionality”

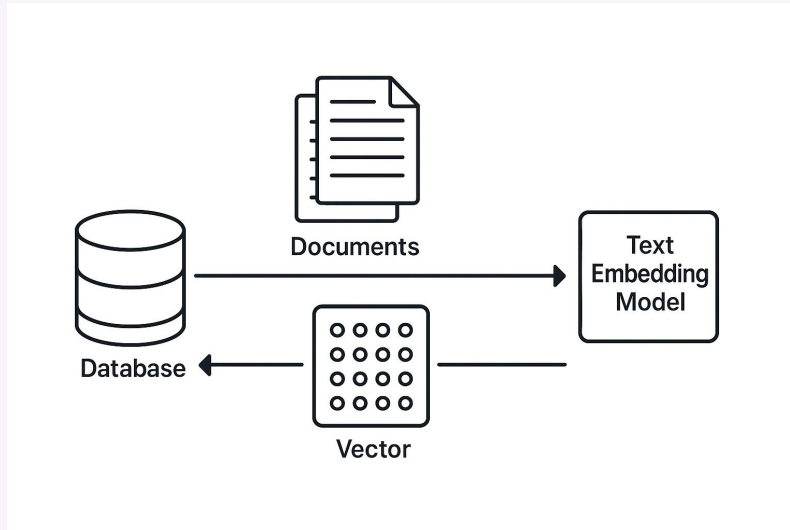
Word2Vec (Google News)	- 300	- one of earliest, not context aware
FastText (Facebook)	- 300	- good for misspelled words
Universal Sentence Encoder (Google)	- 512	- semantic tasks and question answering
BERT	- 768	
OpenAI Ada v2	- 1536	

- We have no idea what attributes are :)
- NOT LLMs!



Vector Search in RavenDB

- Available in v7
- Embedded bge-micro-v2 (384) for fast prototyping
- AI Tasks - Embeddings Generation - support for external models



Secondary aspects

- Caching
 - Reduce latency
 - Reduce cost
 - OpenAI text-embedding-3-large : \$0.00013 / 1k tokens
- Compression
 - OpenAI text-embedding-3-small - 1536
 - 100M embeddings - 572 GB
 - 250M embeddings - 1.430 GB
 - OpenAI text-embedding-3-large - 3072
 - 100M embeddings - 1.144 GB
 - 250M embeddings - 2.861 GB
 - Quantization
 - Single (no quantization)
 - Int8 (float32 -> int8, scaling)
 - fast, moderate accuracy loss, 4x smaller
 - Binary (float32 -> 1 binary bit) [0.2, -0.5, 1.3] -> [1, 0, 1] (threshold is 0)
 - exploits overparameterization, use on larger models, e.g. over 1024
 - extremely fast, high accuracy loss, 32x smaller



Not just search!

- Clustering
 - Group similar items together
 - Detect emerging trends
 - Predict missing relationships
- Recommendation systems
 - Find similar items
- Anomaly detection
 - Vectors far from normal distribution
- Content deduplication
 - Detect nearly-identical items
- User Profiling / Behavior Modelling

